

**Fazang: A Reverse-mode
Automatic differentiation tool in
Fortran**

**User's Guide
(Version 0.0.1)**

Yi Zhang

Contents

Chapter 1. Introduction	4
Chapter 2. Quick Start	5
Chapter 3. Use Fazang	7
1. Constructors	7
2. Assignment	7
3. Gradient	7
4. Functions	9
Chapter 4. Design	10
1. tape data structure	10
2. vari type	10
3. var type	11
Chapter 5. Add operation functions	13
Appendix A. Fazang Functions	16
Appendix. Bibliography	17

CHAPTER 1

Introduction

Fazang is a reverse-mode automatic differentiation (AD) tool. The project is heavily influenced by **Stan/Math** [1], a project the author is also involved in. **Fazang** is intended to support general scientific computing in Fortran beyond Bayesian inference and Markove Chain Monte Carlo that **Stan/Math** is designed for.

User should be aware that the project is at early stage and still under development. For any questions, suggestions, and contributions, please visit the project at <https://github.com/yizhang-yiz/fazang>.

CHAPTER 2

Quick Start

Currently **Fazang** has been tested on Linux and MacOS platform, with Fortran compiler Intel Fortran 19.0.1+ and GNU Fortran 11.2.0+.

After downloading **Fazang**, user can use **meson** to build the library.

```
git clone git@github.com:yizhang-yiz/fazang.git
cd fazang && mkdir build && cd build
meson compile
```

This generates a shared library at **build/src/**. User needs to link this library when building an application. This can be done in **meson** by setting

```
executable('app_name', files('path/to/app_file.F90'),
↳ dependencies : fazang_dep)
```

Fazang provides a user-facing derived type **var**. This is the type for the dependent and independent variables of which the adjoint (derivative) will be calculated.

For example, consider the log of the Gaussian distribution density with mean μ and standard deviation σ

$$f(\mu, \sigma) = \log \left(\frac{1}{\sigma \sqrt{2\pi}} \exp \left(-\frac{1}{2} \left(\frac{y - \mu}{\sigma} \right)^2 \right) \right) \quad (1)$$

The following programme calculates $\frac{df}{d\mu}$ and $\frac{df}{d\sigma}$ at $y = 1.3$, $\mu = 0.5$, and $\sigma = 1.2$.

```
program log_demo
  use fazang ! load Fazang library

  implicit none

  real(rk) :: y
```

```
type(var) :: f, sigma, mu

! data
y = 1.3d0

! independent variables
mu = var(0.5d0)
sigma = var(1.2d0)

! dependent
f = var(-0.5d0 * log(2 * pi))
f = f - log(sigma)
f = f - 0.5d0 * ((y - mu) / sigma) ** 2.d0;

! use grad() to calculate df/d(mu) and df/d(sigma). Each var's
! derivative (also called adjoint) can be access through
! ↪ var%adj().

call f%grad()
write(*, *) "df/d(mu): ", mu%adj()
write(*, *) "df/d(sigma): ", sigma%adj()
end program log_demo
```

CHAPTER 3

Use Fazang

Fazang uses `var` type to record numerical and gradient operations. The type supports three functions

- `var%val()` : returns value
- `var%adj()` : returns derivative, henceforth referred as *adjoint*.
- `var%grad()` : takes gradient operation with respect to the current `var` variable.

1. Constructors

`var` can be constructed using overloaded `var` interface.

```
real(real64) :: a, b(3), c(2, 3)
real(real64) :: new_a, new_b(3), new_c(2, 3)
type(var) :: x, y(3), z(2, 3)
! ...
x = var()                ! x%val() == 0.d0
x = var(a)               ! x%val() == a
y = var(b)               ! y%val() == b
z = var(c)               ! z%val() == c
```

2. Assignment

`var` can be assigned from consistent `var` and `real(real64)`.

```
! ....
x = new_a                ! x%val() == new_a
y = new_b                ! y%val() == new_b
z = new_c                ! z%val() == new_c
```

3. Gradient

Consider a variable z calculated by the composition of a series of operations

$$z = f_1(z_1), \quad z_1 = f_2(z_2), \quad \dots, \quad z_{n-1} = f_n(z_n).$$

For $z_i, i = 1, \dots, n$ we refer dz/dz_i as the *adjoint* of z_i , denoted by z_i^{adj} . The chain rule says the adjoints can be calculated recursively [2],

$$\begin{aligned} z^{\text{adj}} &= 1, \\ z_1^{\text{adj}} &= z^{\text{adj}} \frac{df_1}{dz_1}, \\ &\dots, \\ z_i^{\text{adj}} &= z_{i-1}^{\text{adj}} \frac{df_i}{dz_i}. \end{aligned}$$

We often refer each (f_i, z_i) pair as a *node*, and z_i the *operand* of operation f_i . The above recursion through the nodes requires a way to store and visit the *callstack* of nodes. It is embodied in **Fazang** by the `var%grad()` function. When `z%grad()` is called, `z`'s adjoint is set to 1, and every other `var` variable is transversed with its adjoint updated. In order to calculate the adjoint with respect to another variable, user must call `set_zero_all_adj()` first to reset all adjoints to zero.

An alternative to invoke gradient calculation is to define the dependent as a function and feed it to **Fazang**'s `gradient` function. Take Eq.(1) for example, we can first define the function for $f(\mu, \sigma)$.

```
module func
  use fazang ! load Fazang library
  implicit none

  real(rk), parameter :: y = 1.3d0

  contains
    type(var) function f(x)
      type(var), intent(in) :: x(:)
      type(var) :: mu, sigma
      mu = x(1)
      sigma = x(2)
      f = -0.5d0 * log(2 * pi) - log(sigma) - 0.5d0 * ((y - mu) /
        ↪ sigma) ** 2.d0;
    end function f
end module func
```

Then we can supply function `f` as a procedure argument.

```
program log_demo2
  use iso_c_binding
```



```
use fazang
use func

implicit none

real(real64) :: fx(3), x(2)
x = [0.5d0, 1.2d0]

fx = gradient(f, x)
write(*, *) "f(x): ", fx(1)
write(*, *) "df/d(x(1)):", fx(2)
write(*, *) "df/d(x(2)):", fx(3)
end program log_demo2
```

The output of `gradient(f, x)` is an array of size `1 + size(x)`, with first component being the function value, and the rest the partial derivatives.

4. Functions

Numeric functions supported by **Fazang** are listed in Appendix A. All unary and binary functions are **elemental**. The binary functions allow mixed argument types, namely, either argument can be **real64** type while the other the **var** type.

CHAPTER 4

Design

The core of any reverse-mode automatic differentiation is the data structure to store and visit the callstack. **Fazang** achieves this through two derived types, **tape** and **vari**.

1. tape data structure

A **tape** is an **int32** array emulating a stack, with an integer marker **head** pointing to the head to the current stack top.

```
type :: tape
  integer(ik) :: head = 1
  integer(ik), allocatable :: storage(:)
  !...
```

Each time a new AD node is created, space in **storage** is allotted to store the node's

- value $f_i(z_i)$,
- adjoint z_{i-1}^{adj} ,
- number of **var** operands of f_i ,
- The **var** operands' index in the same **tape** array,
- number of **real64** operands of f_i ,
- The **real64** operands' value.

Since a node's value, adjoint, and data operands are **real64**, they are first converted to **int32** using **transfer** function before stored in the **tape** array, so that each such a value occupies two **storage** entries. After each allotation, the **head** is moved to point to the next empty slot in the array after saving its current value to a **vari** type variable for future retrieval.

2. vari type

The **vari** type is simply a proxy of a node's storage location in the **tape**

```

type :: vari
  integer(ik) :: i = 0
  procedure(chain_op), pass, pointer :: chain
contains
  !....

```

where *i* is the index to the beginning of a node's storage, and the `chain` procedure encodes the node's operation f_i . `chain` follows an interface that describes the chain rule operation

```

abstract interface
  subroutine chain_op(this)
    import :: vari
    class(vari), intent(in) :: this
  end subroutine chain_op
end interface

```

An alternative to integer index is to a `pointer` to the according entry in the `tape` array. However, we will need to expand the `storage` when it is filled up, and `Fazang` does this by doubling the `storage` size and use `move_alloc` to restore the original values. Since there is no guarantee that `move_alloc` will keep the original memory, a pointer to the original address would be corrupted.

As a `Fazang` program steps forward, a series of `vari` variables are generated, with their *values* calculated and stored. This is called a *forward pass*. The generated `vari` variables in the forward pass are stored in array `varis`. Each entry in `varis` is a dependent (operation output) of one or more previous entries.

3. var type

The user-facing `var` type serves as proxy to `vari`. Each `var` stores the index of a `vari` in the `varis` array.

```

type :: var
  integer(int32) :: vi
contains
  procedure :: val
  procedure :: adj
  procedure :: grad
  procedure :: set_chain
end type var

```

After the forward pass, when adjoints are desired, we call `grad` or `gradient` procedure. This initiates a *reverse pass*, in which the `varis` array is transversed backward so that each `vari`'s `chain` procedure is called to update the operand adjoints.

```
subroutine grad(this)
  class(var), intent(in) :: this
  integer i
  call callstack % varis (this%vi) % init_dependent()
  do i = callstack % head - 1, 1, -1
    call callstack % varis(i) % chain()
  end do
end subroutine grad
```

Here `callstack` is the module variable that encapsulate `tape` and `varis` arrays.

CHAPTER 5

Add operation functions

Adding an operation f_i involves creating functions for forward pass and reverse pass. Let us first use `log` function as a simple example.

First, we create a `log_v` function for the forward pass.

```
impure elemental function log_v(v) result(s)
  type(var), intent(in) :: v
  type(var) :: s
  s = var(log(v%val()), [v])
  call s%set_chain(chain_log)
end function log_v
```

The function generates a new `var` variable `s` using a special constructor `var(value, array of operands)` which stores the value as well as the single operand `v`'s index (in the `tape storage` array). It also points `s`'s chain to a dedicated procedure `chain_log`.

```
subroutine chain_log(this)
  class(vari), intent(in) :: this
  real(rk) :: new_adj(1), val(1)
  new_adj = this%operand_adj()
  val = this%operand_val()
  new_adj(1) = new_adj(1) + this%adj() / val(1)
  call this%set_operand_adj(new_adj)
end subroutine chain_log
```

To understand this function, recall the recursion in Section 3, assume the `log` operation is node i , then $f_i = \log()$ and z_i is the operand `v`, and the new `var s` would be z_{i-1} . During the reverse pass when the node is visited, `chain_log` first retrieves current (z_i, z_i^{adj}) using `operand_val()` and `operand_adj()`, then updates z_i^{adj} with an additional

$$z_{i-1}^{\text{adj}} \frac{df_i}{dz_i} = z_{i-1}^{\text{adj}} \frac{d \log(z_i)}{dz_i} = \frac{z_{i-1}^{\text{adj}}}{z_i}.$$

Adding a binary operation $f_i(z_i^{(1)}, z_i^{(2)})$ is slightly more complex, as we will need to address possibly different scenarios when $z_i^{(1)}$ and $z_i^{(2)}$ are either `var` or `real64`. Let us use overloaded division operator `(/)` as an example.

With

$$f_i(z_i^{(1)}, z_i^{(2)}) = z_i^{(1)} / z_i^{(2)}$$

we need to account for

- both $z_i^{(1)}$ and $z_i^{(2)}$ are `var`'s
- $z_i^{(1)}$ is `var`, $z_i^{(2)}$ is `real64`,
- $z_i^{(1)}$ is `real64`, $z_i^{(2)}$ is `var`,

For the first scenario, we create

```
impure elemental function div_vv(v1, v2) result(s)
  type(var), intent(in) :: v1, v2
  type(var) :: s
  s = var(v1%val() / v2%val(), [v1, v2])
  call s%set_chain(chain_div_vv)
end function div_vv
```

Similar to the `log` example, we create a new `s` with both operands stored. In the corresponding `chain` procedure, we need update the adjoints of both `v1` and `v2`.

```
subroutine chain_div_vv(this)
  class(vari), intent(in) :: this
  real(rk) :: new_adj(2), val(2)
  new_adj = this%operand_adj()
  val = this%operand_val()
  new_adj(1) = new_adj(1) + this%adj()/val(2)
  new_adj(2) = new_adj(2) - this%val() * this%adj()/val(2)
  call this%set_operand_adj(new_adj)
end subroutine chain_div_vv
```

For the second scenario, we create

```
impure elemental function div_vd(v, d) result(s)
  type(var), intent(in) :: v
  real(rk), intent(in) :: d
  type(var) :: s
  s = var(v%val() / d, [v], [d])
  call s%set_chain(chain_div_vd)
end function div_vd
```

Again we create a new `var` `s`. But this time we use another constructor `var(value, var operands, data operands)` to store value, `var` operand `v`, and `real64` operand `d`. In the corresponding reverse pass `chain` procedure, not only we need retrieve `var` operand `v` but also data operand `d`, as the new adjoint of $z_i^{(1)}$ is

$$z_i^{(1)\text{new adj}} = z_i^{(1)\text{old adj}} + z_{i-1}^{\text{adj}} \frac{df_i}{dz_i^{(1)}} = z_i^{(1)\text{old adj}} + z_{i-1}^{\text{adj}} \frac{1}{dz_i^{(2)}}$$

So with `v` as $z_i^{(1)}$ and `d` as $z_i^{(2)}$ we have

```
subroutine chain_div_vd(this)
  class(vari), intent(in) :: this
  real(rk) d(1), new_adj(1)
  new_adj = this%operand_adj()
  d = this%data_operand()
  new_adj(1) = new_adj(1) + this%adj() / d(1)
  call this%set_operand_adj(new_adj)
end subroutine chain_div_vd
```

The third scenario is treated similarly.

APPENDIX A

Fazang Functions

Function	Argument(s)	Operation
<code>sin</code>	scalar or array	same as intrinsic
<code>cos</code>	scalar or array	same as intrinsic
<code>tan</code>	scalar or array	same as intrinsic
<code>asin</code>	scalar or array	same as intrinsic
<code>acos</code>	scalar or array	same as intrinsic
<code>atan</code>	scalar or array	same as intrinsic
<code>log</code>	scalar or array	same as intrinsic
<code>exp</code>	scalar or array	same as intrinsic
<code>sqrt</code>	scalar or array	same as intrinsic
<code>square</code>	scalar or array	For input x , calculate $x**2$
<code>inv</code>	scalar or array	For input x , calculate $1/x$
<code>inv_square</code>	scalar or array	For input x , calculate $1/x**2$
<code>inv_sqrt</code>	scalar or array	For input x , calculate $1/\sqrt{x}$
<code>logit</code>	scalar or array	For input x , calculate $\log(x/(1-x))$
<code>inv_logit</code>	scalar or array	For input x , calculate $1/(1+\exp(-x))$
operator (+)	scalars or arrays	same as intrinsic
operator (-)	scalars or arrays	same as intrinsic
operator (*)	scalars or arrays	same as intrinsic
operator (/)	scalars or arrays	same as intrinsic
operator (**)	scalars	same as intrinsic
<code>sum</code>	1D array	same as intrinsic
<code>dot_product</code>	1D arrays	same as intrinsic
<code>log_sum_exp</code>	1D array	For input x , calculate $\log(\sum(\exp((x))))$
<code>matmul</code>	2D arrays	same as intrinsic

Bibliography

- [1] Bob Carpenter, Matthew D. Hoffman, Marcus A. Brubaker, Daniel Lee, Peter Li, and Michael J. Betancourt. The Stan math library: Reverse-mode automatic differentiation in C++. *arXiv 1509.07164.*, 2015. <https://mc-stan.org/users/interfaces/math>.
- [2] Andreas Griewank and Andrea Walther. *Evaluating Derivatives: Principles and Techniques of Algorithmic Differentiation, Second Edition*. Society for Industrial and Applied Mathematics, Philadelphia, PA, second edition edition, September 2008.